

Async Benchmarks Index

Mar 22, 2021

I don't understand performance characteristics of “async” programming when applied to typical HTTP based web applications. Let's say we have a CRUD app with a relational database, where a typical request results in N queries to the database and transfers M bytes over the network. How much (orders of magnitude?) faster/slower would an “async” solution be in comparison to a “threaded” solution?

In this *live* post, I am collecting the benchmarks that help to shed the light on this and related questions. Note that I am definitely not the right person to do this work, so, if there is a better resource, I'll gladly just use that instead. Feel free to [send pull requests](#) with benchmarks! Every benchmark will be added, but some might go to the rejected section.

I am interested in understanding differences between several execution models, regardless of programming language:

Threads

Good old POSIX threads, as implemented on modern Linux.

Stackful Coroutines

M:N threading, which expose the same programming model as threads, but are implemented by multiplexing several user-space coroutines over a single OS-level thread. The most prominent example here is Go

Stackless Coroutines

In this model, each concurrent computation is represented by a fixed-size state machine which reacts to events. This model often uses `async / await` syntax for describing and composing state machines using standard control

flow constructs.

Threads With Cooperative Scheduling

This is a mostly hypothetical model of OS threads with an additional primitive for directly switching between two threads of the same process. It is not implemented on Linux (see [this presentation](#) for some old work towards that). It is implemented on Windows under the “fiber” branding.

I am also interested in Rust’s specific implementation of stackless coroutines

Benchmarks

<https://github.com/jimblandy/context-switch>

This is a micro benchmark comparing the cost of primitive operations of threads and stackless as implemented in Rust coroutines. Findings:

- Thread creation is order of magnitude slower
- Threads use order of magnitude more RAM.
- IO-related context switches take the same time
- Thread-to-thread context switches (channel sends) take the same time, *if* threads are pinned to one core. This is surprising to me. I’d expect channel send to be significantly more efficient for either stackful or stackless coroutines.
- Thread-to-thread context switches are order of magnitude slower if there’s no pinning
- Threads hit non-memory resource limitations quickly (it’s hard to spawn > 50k threads).

<https://github.com/jkarneges/rust-async-bench>

Micro benchmark which compares Rust’s implementation of stackless corou-

tines with a manually coded state machine. Rust's async/await turns out to not be zero-cost, pure overhead is about 4x. The absolute numbers are still low though, and adding even a single syscall of work reduces the difference to only 10%

<https://matklad.github.io/2021/03/12/goroutines-are-not-significantly-smaller-than-threads.html>

This is a micro benchmark comparing just the memory overhead of threads and stackful coroutines as implemented in Go. Threads are “times”, but not “orders of magnitude” larger.

<https://calpaterson.com/async-python-is-not-faster.html>

Macro benchmark which compares many different Python web frameworks. The conclusion is that async is worse for both latency and throughput. Note two important things. *First*, the servers are run behind a reverse proxy (nginx), which drastically changes IO patterns that are observed by the server. *Second*, Python is not the fastest language, so throughput is roughly correlated with the amount of C code in the stack.

There is also [a rebuttal post](#).

Rejected Benchmarks

<https://matej.laitl.cz/bench-actix-rocket/>

This is a macro benchmark comparing performance of sync and async Rust web servers. This is the kind of benchmark I want to see, and the analysis is exceptionally good. Sadly, a big part of the analysis is fighting with unreleased version of software and working around bugs, so I don't trust that the results are representative.

<https://www.techempower.com/benchmarks/>

This is a micro benchmark that pretends to be a macro benchmark. The code

is overly optimized to fit a very specific task. I don't think the results are easily transferable to real-world applications. At the same time, lack of the analysis and the "macro" scale of the task itself doesn't help with building a mental model for explaining the observed performance.

<https://inside.java/2020/08/07/loom-performance>

The opposite of a benchmark actually. This post gives a good theoretical overview of why async might lead to performance improvements. Sadly, it drops the ball when it comes to practice:

millions of user-mode threads instead of the meager thousands the OS can support.

What is the limiting factor for OS threads?